

# Employee Wellness, Performance, and Retention Analysis

```
In [1]: import os;
os.listdir ('/kaggle/input/')

Out[1]: ['wellness2', 'wellness']
```

## IMPORTING THE DATA SET

```
In [2]: #import Library
import pandas as pd

#Read each sheet into a Dataframe

df_sheet1 = pd.read_excel('/kaggle/input/wellness/Vacation and wellness.xlsx', sheet_name = 'Vacation n Wellness 2022')
df_sheet2 = pd.read_excel('/kaggle/input/wellness/Vacation and wellness.xlsx', sheet_name = 'Vacation n Wellness 2023')
```

## Merged sheets and Display the first few rows

```
In [3]: #Merged the sheets using a full outer join

merged_df = pd.merge (df_sheet1, df_sheet2, on = 'Employee ID', how = 'outer')

#Display the first few rows of the merged Dataframe

merged_df.head()
```

Out[3]:

	Employee ID	Gender_x	Hire Date_x	Position Title_x	Grade_x	Division Classification_x	Company_x	Career Stream_x	Event Date_x	Employee Status_x	...	Event	Term
0	24	Woman	2003-05-26	Mgr Client Services & Administration	P7	MI-ISS	Mackenzie Financial Corporation	Professional	2020-02-19	Terminated	...	Termination	202
1	101	Woman	1982-08-09	Snr Mgr Client Act Mgt & Special Sup	L8	IGM-CSO	Investors Group Financial Services	Leader	2022-10-16	Active	...	Data Change	
2	104	Woman	1982-12-13	Loss Mitigation Analyst	P5	IG-InsMtgBnk	Investors Group Investment Management	Professional	2022-03-01	Active	...	Data Change	202
3	161	Man	1980-10-27	Director Real Estate Asset Mgmt	P9	IG-IGIM	Investors Group Investment Management	Professional	2021-01-01	Terminated	...	Termination	202
4	311	Woman	1977-07-04	Snr Mgr Learning & Communication	L8	IGM-CSO	Investors Group Financial Services	Leader	2020-02-12	Terminated	...	Termination	202

5 rows × 46 columns

## DISPLAY BASIC INFORMATION AND STATISTICS OF THE DATAFRAME

```
In [4]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Display basic info and statistics of the DataFrame

print(merged_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6929 entries, 0 to 6928
Data columns (total 46 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Employee ID                               6929 non-null   int64
1   Gender_x                                  6044 non-null   object
2   Hire Date_x                               6044 non-null   datetime64[ns]
3   Position Title_x                         6043 non-null   object
4   Grade_x                                   6044 non-null   object
5   Division Classification_x                 6044 non-null   object
6   Company_x                                6044 non-null   object
7   Career Stream_x                          6044 non-null   object
8   Event Date_x                             6044 non-null   datetime64[ns]
9   Employee Status_x                        6044 non-null   object
10  Termination Date_x                       1968 non-null   datetime64[ns]
11  STIP Total Target_x                       5747 non-null   object
12  STIP Total Max_x                          5747 non-null   object
13  Vacation Annual Entitlement_x             6043 non-null   float64
14  Age_x                                     6044 non-null   float64
15  Wellness - Used_x                         6044 non-null   float64
16  Wellness - Accrued_x                     6044 non-null   float64
17  Wellness - %                             6044 non-null   float64
18  Wellness Used round up                    6044 non-null   object
19  Vacation - Used_x                         6044 non-null   float64
20  Vacation - Accrued_x                     6044 non-null   float64
21  Vacation - %                             6044 non-null   float64
22  Vacation Used round up                    6044 non-null   object
23  Age_y                                     6916 non-null   float64
24  Gender_y                                  6916 non-null   object
25  Employee Status_y                        6916 non-null   object
26  Company_y                                6916 non-null   object
27  Division Classification_y                 6916 non-null   object
28  Hire Date_y                              6916 non-null   datetime64[ns]
29  Position Title_y                         6916 non-null   object
30  Grade_y                                   6916 non-null   object
31  Career Stream_y                          6916 non-null   object
32  STIP Total Target_y                       6655 non-null   object
33  STIP Total Max_y                          6655 non-null   object
34  Vacation Annual Entitlement_y             6915 non-null   float64
35  Event Date_y                             6916 non-null   datetime64[ns]
36  Event                                     6916 non-null   object
37  Termination Date_y                       2759 non-null   datetime64[ns]
38  Wellness - Used_y                         6916 non-null   float64
39  Wellness - Accrued_y                     6916 non-null   float64
40  Wellness Used- %                          6916 non-null   float64
41  Wellness-round up                         6916 non-null   object
42  Vacation - Used_y                         6916 non-null   float64
43  Vacation - Accrued_y                     6916 non-null   float64
44  Vacation Used - %                         6916 non-null   float64
45  Vacation-round up                         6916 non-null   object
dtypes: datetime64[ns](6), float64(16), int64(1), object(23)
memory usage: 2.4+ MB
None
```

```
In [5]: # Check for missing values
print(merged_df.isnull().sum())
```

Employee ID	0
Gender_x	885
Hire Date_x	885
Position Title_x	886
Grade_x	885
Division Classification_x	885
Company_x	885
Career Stream_x	885
Event Date_x	885
Employee Status_x	885
Termination Date_x	4961
STIP Total Target_x	1182
STIP Total Max_x	1182
Vacation Annual Entitlement_x	886
Age_x	885
Wellness - Used_x	885
Wellness - Accrued_x	885
Wellness - %	885
Wellness Used round up	885
Vacation - Used_x	885
Vacation - Accrued_x	885
Vacation - %	885
Vacation Used round up	885
Age_y	13
Gender_y	13
Employee Status_y	13
Company_y	13
Division Classification_y	13
Hire Date_y	13
Position Title_y	13
Grade_y	13
Career Stream_y	13
STIP Total Target_y	274
STIP Total Max_y	274
Vacation Annual Entitlement_y	14
Event Date_y	13
Event	13
Termination Date_y	4170
Wellness - Used_y	13
Wellness - Accrued_y	13
Wellness Used- %	13
Wellness-round up	13
Vacation - Used_y	13
Vacation - Accrued_y	13
Vacation Used - %	13
Vacation-round up	13

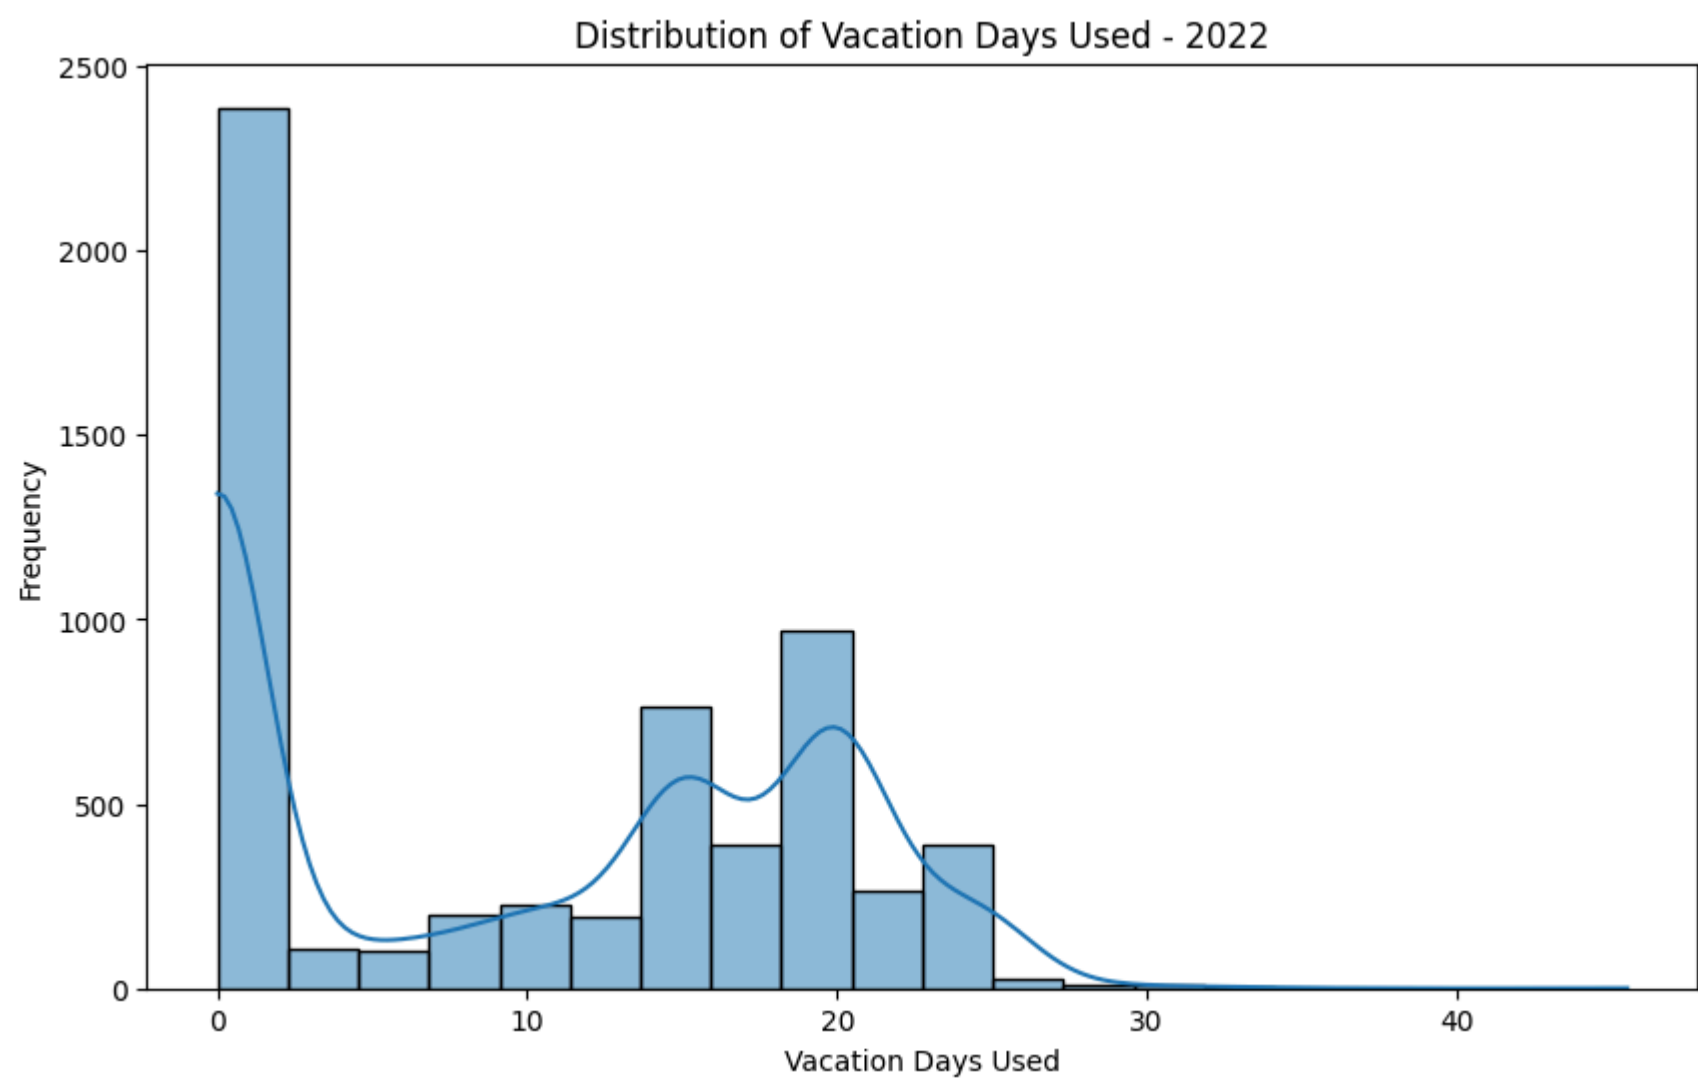
dtype: int64

## DISTRIBUTION OF VACATION DAYS USED

```
In [6]: import matplotlib.pyplot as plt
import seaborn as sns


# Plot distribution of vacation days used for 2022
plt.figure(figsize=(10, 6))
sns.histplot(merged_df['Vacation - Used_x'], bins=20, kde=True)
plt.title('Distribution of Vacation Days Used - 2022')
plt.xlabel('Vacation Days Used')
plt.ylabel('Frequency')
plt.show()
```

/opt/conda/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):



```
In [7]: # Plot distribution of vacation days used for 2023
plt.figure(figsize=(10, 6))
sns.histplot(merged_df['Vacation - Used_y'], bins=20, kde=True)
plt.title('Distribution of Vacation Days Used - 2023')
plt.xlabel('Vacation Days Used')
plt.ylabel('Frequency')
plt.show()
```

/opt/conda/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):




# CORRELATION MATRIX

```
In [8]: # Get numeric columns for correlation matrix
numeric_columns = merged_df.select_dtypes(include=['number']).columns
```

```
In [9]: # Calculate correlation matrix for numeric columns only
corr_matrix = merged_df[numeric_columns].corr()
```

```
In [10]: # Plot correlation matrix as a heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
```

```
plt.title('Correlation Matrix')
plt.show()
```



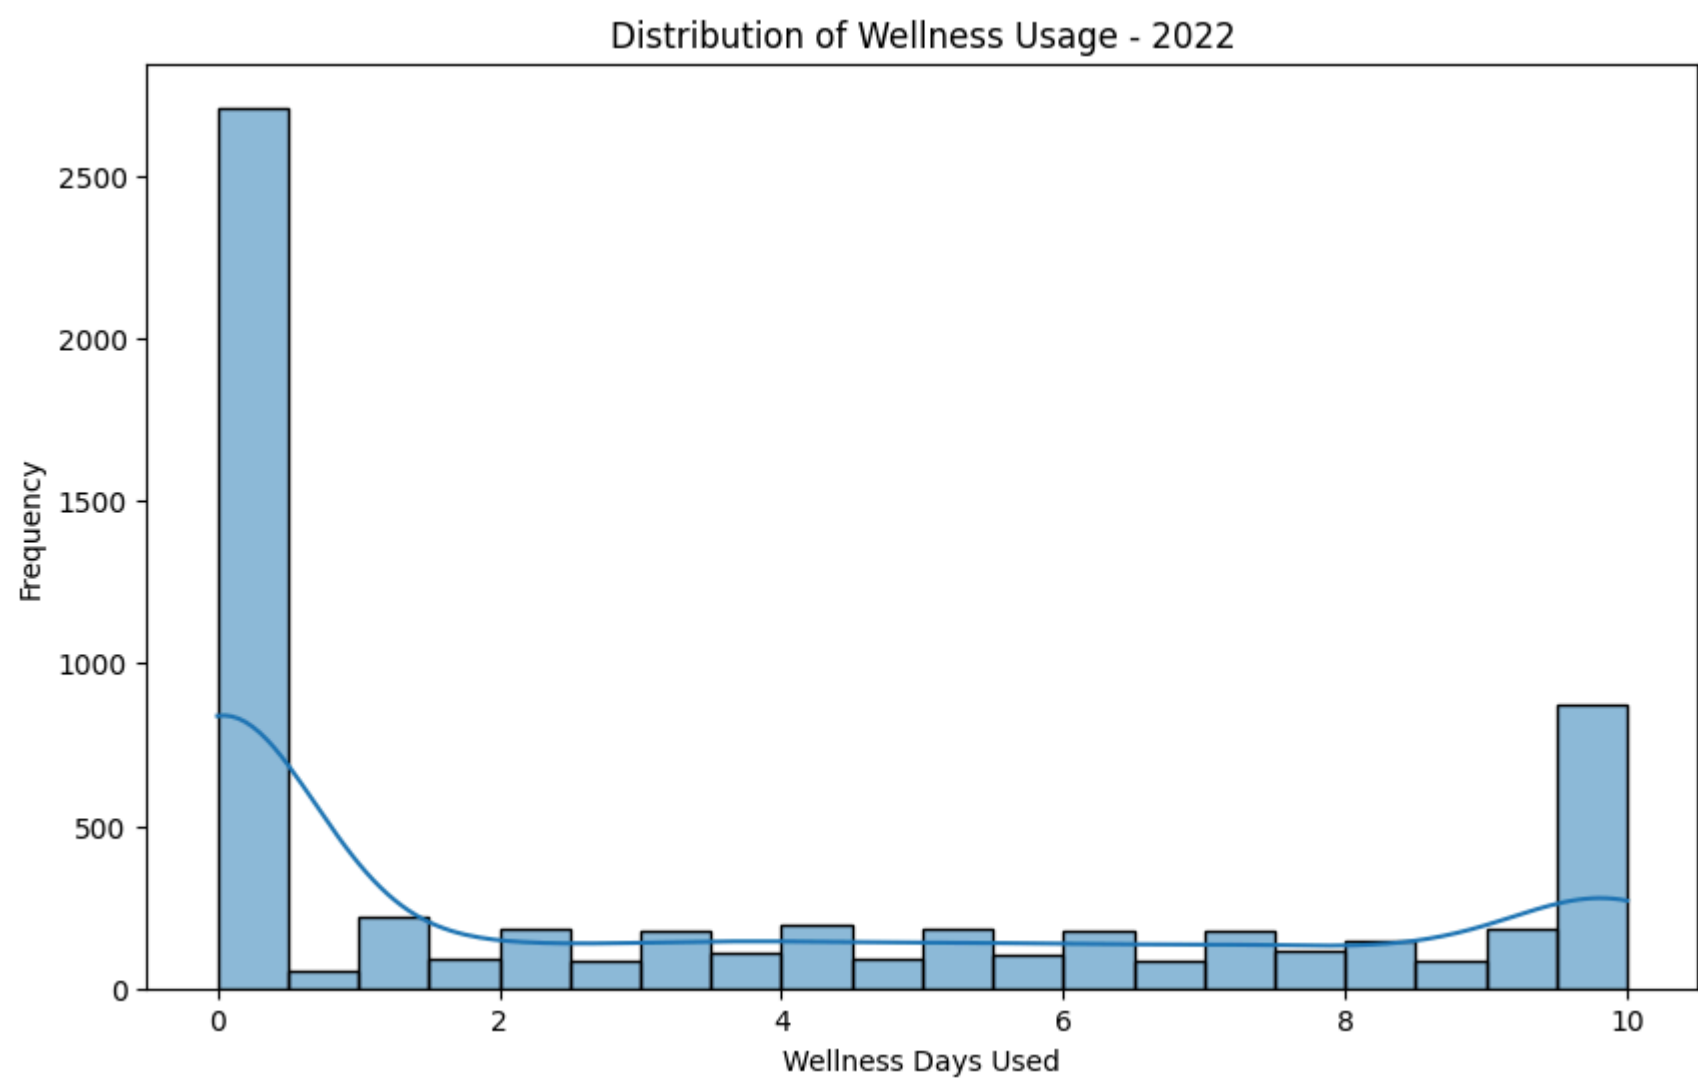
## DISTRIBUTION OF WELLNESS USAGE

```
In [11]: import matplotlib.pyplot as plt
import seaborn as sns

# Plot distribution of wellness usage for 2022
plt.figure(figsize=(10, 6))
sns.histplot(merged_df['Wellness - Used_x'], bins=20, kde=True)
plt.title('Distribution of Wellness Usage - 2022')
plt.xlabel('Wellness Days Used')
plt.ylabel('Frequency')
plt.show()
```

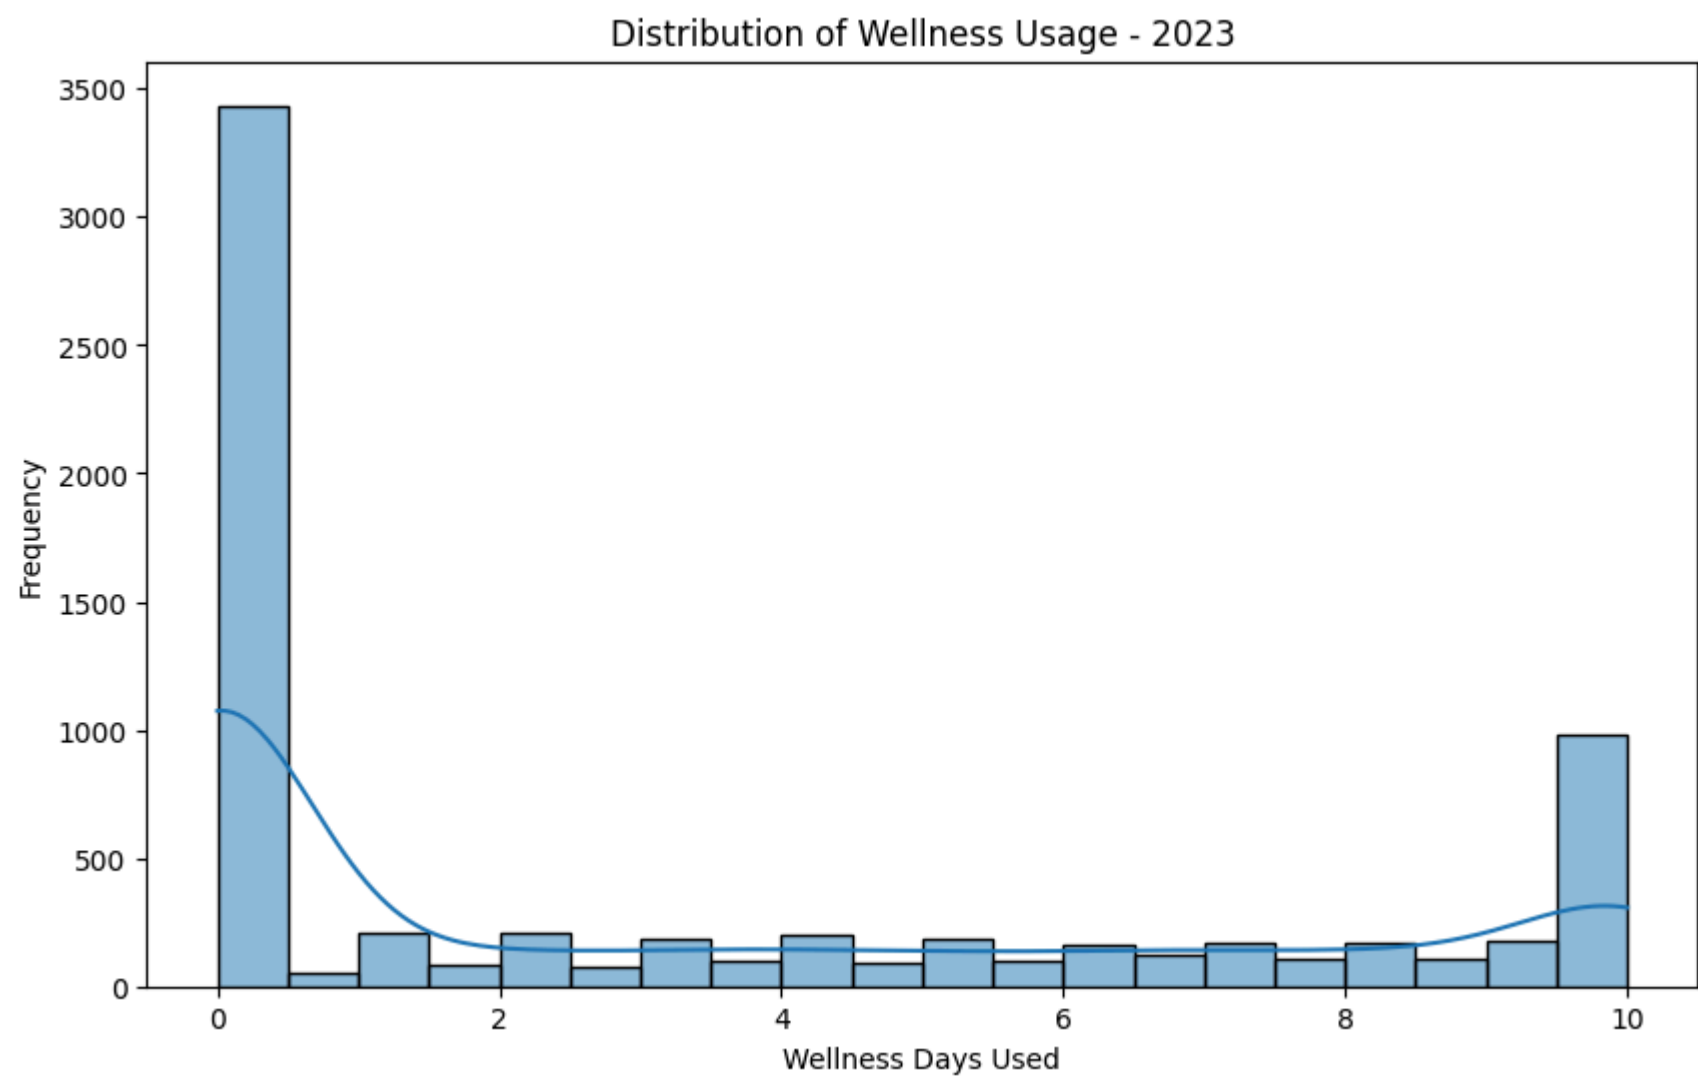
/opt/conda/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

with pd.option\_context('mode.use\_inf\_as\_na', True):



```
In [12]: # Plot distribution of wellness usage for 2023
plt.figure(figsize=(10, 6))
sns.histplot(merged_df['Wellness - Used_y'], bins=20, kde=True)
plt.title('Distribution of Wellness Usage - 2023')
plt.xlabel('Wellness Days Used')
plt.ylabel('Frequency')
plt.show()
```

/opt/conda/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):

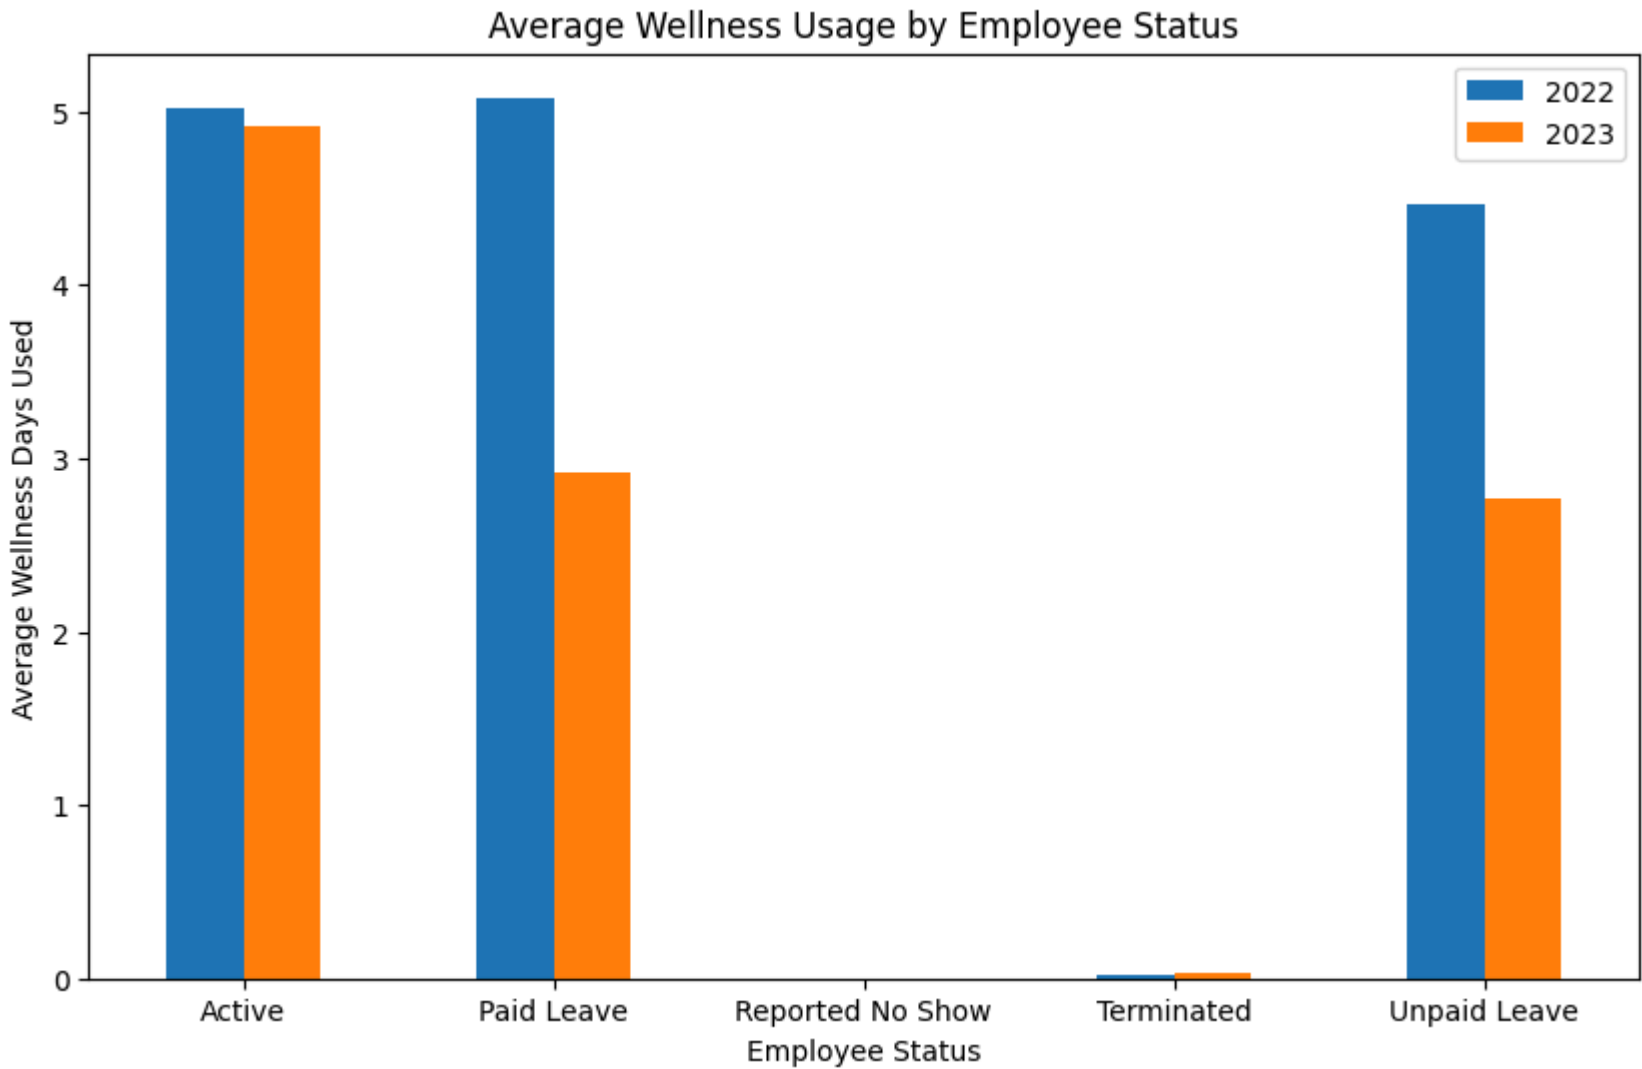


## AVERAGE WELLNESS USED BY EMPLOYEE STATUS

```
In [13]: # Group data by Employee Status and calculate average wellness usage for both years
wellness_grouped = merged_df.groupby('Employee Status_x')[['Wellness - Used_x', 'Wellness - Used_y']].mean()


# Plot average wellness usage by Employee Status
wellness_grouped.plot(kind='bar', figsize=(10, 6))
plt.title('Average Wellness Usage by Employee Status')
plt.xlabel('Employee Status')
plt.ylabel('Average Wellness Days Used')
plt.xticks(rotation=0)
```

```
plt.legend(['2022', '2023'])
plt.show()
```



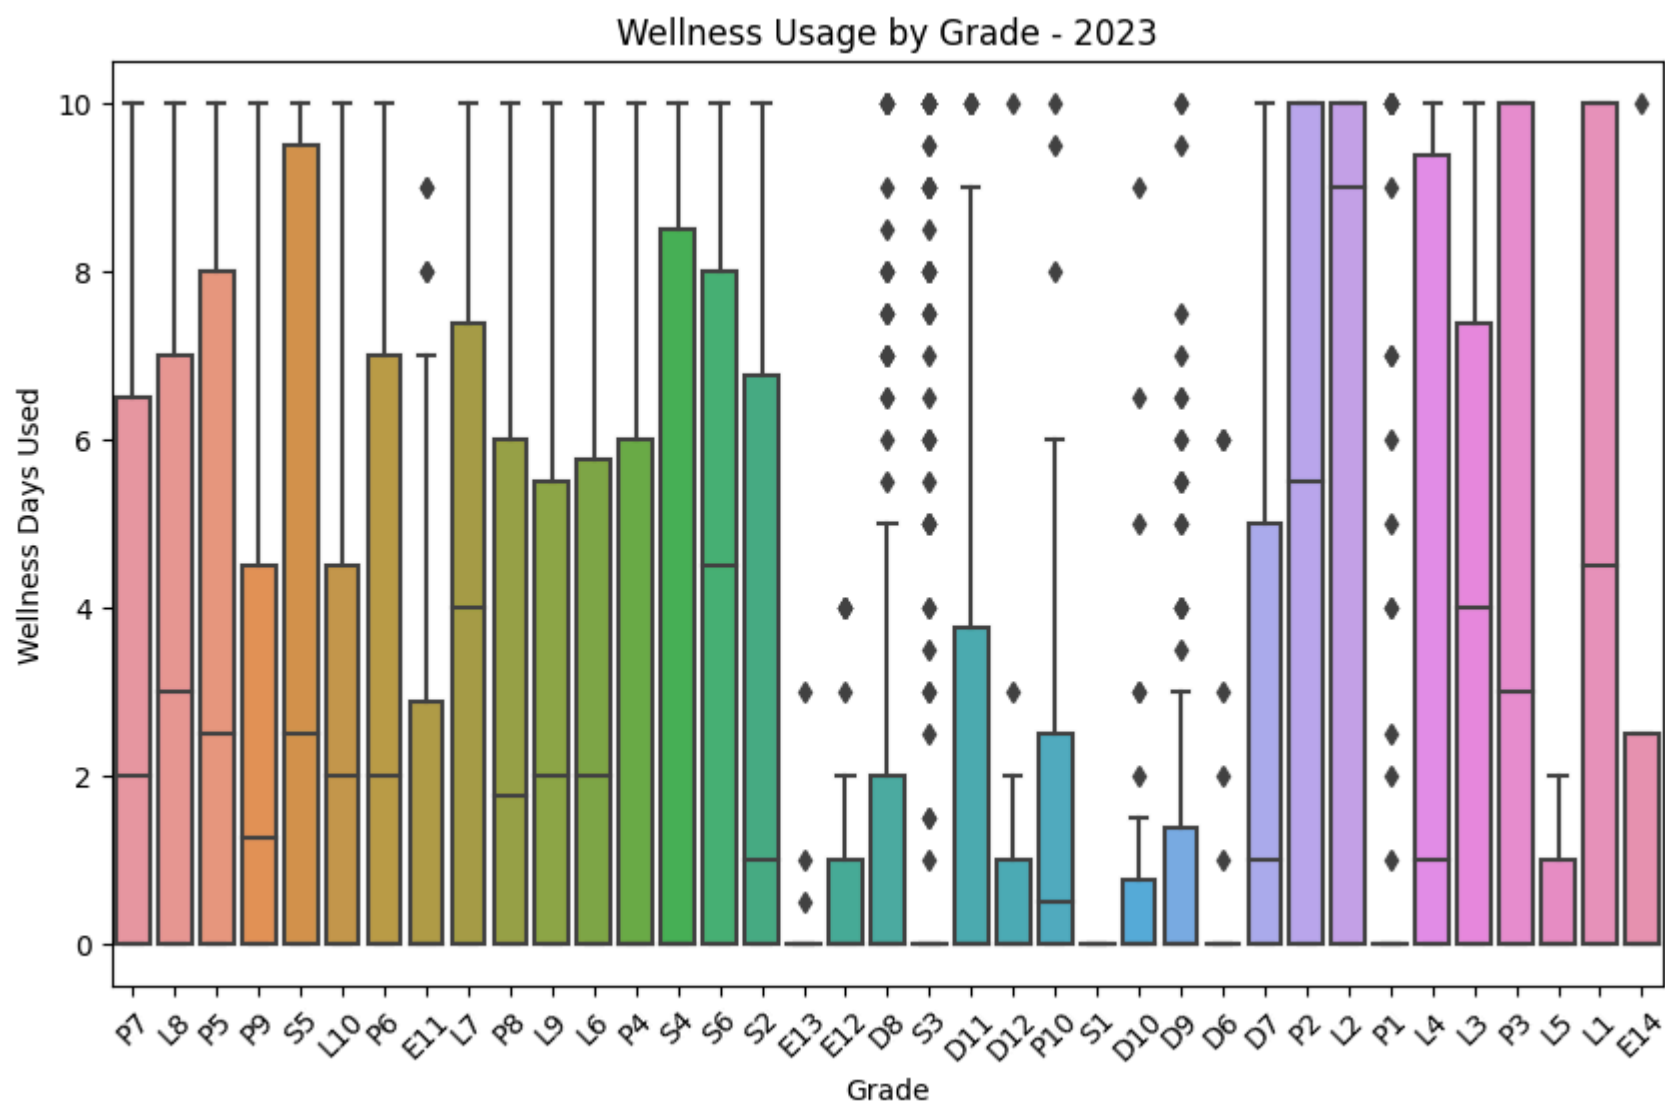
## WELLNESS USAGE BY GRADE

```
In [14]: # Plot box plot of wellness usage by Grade for 2022
plt.figure(figsize=(10, 6))
sns.boxplot(x='Grade_x', y='Wellness - Used_x', data=merged_df)
plt.title('Wellness Usage by Grade - 2022')
plt.xlabel('Grade')
plt.ylabel('Wellness Days Used')
plt.xticks(rotation=45)
plt.show()
```




```
In [15]: # Plot box plot of wellness usage by Grade for 2023
plt.figure(figsize=(10, 6))
sns.boxplot(x='Grade_y', y='Wellness - Used_y', data=merged_df)
plt.title('Wellness Usage by Grade - 2023')
plt.xlabel('Grade')
plt.ylabel('Wellness Days Used')
plt.xticks(rotation=45)
plt.show()
```





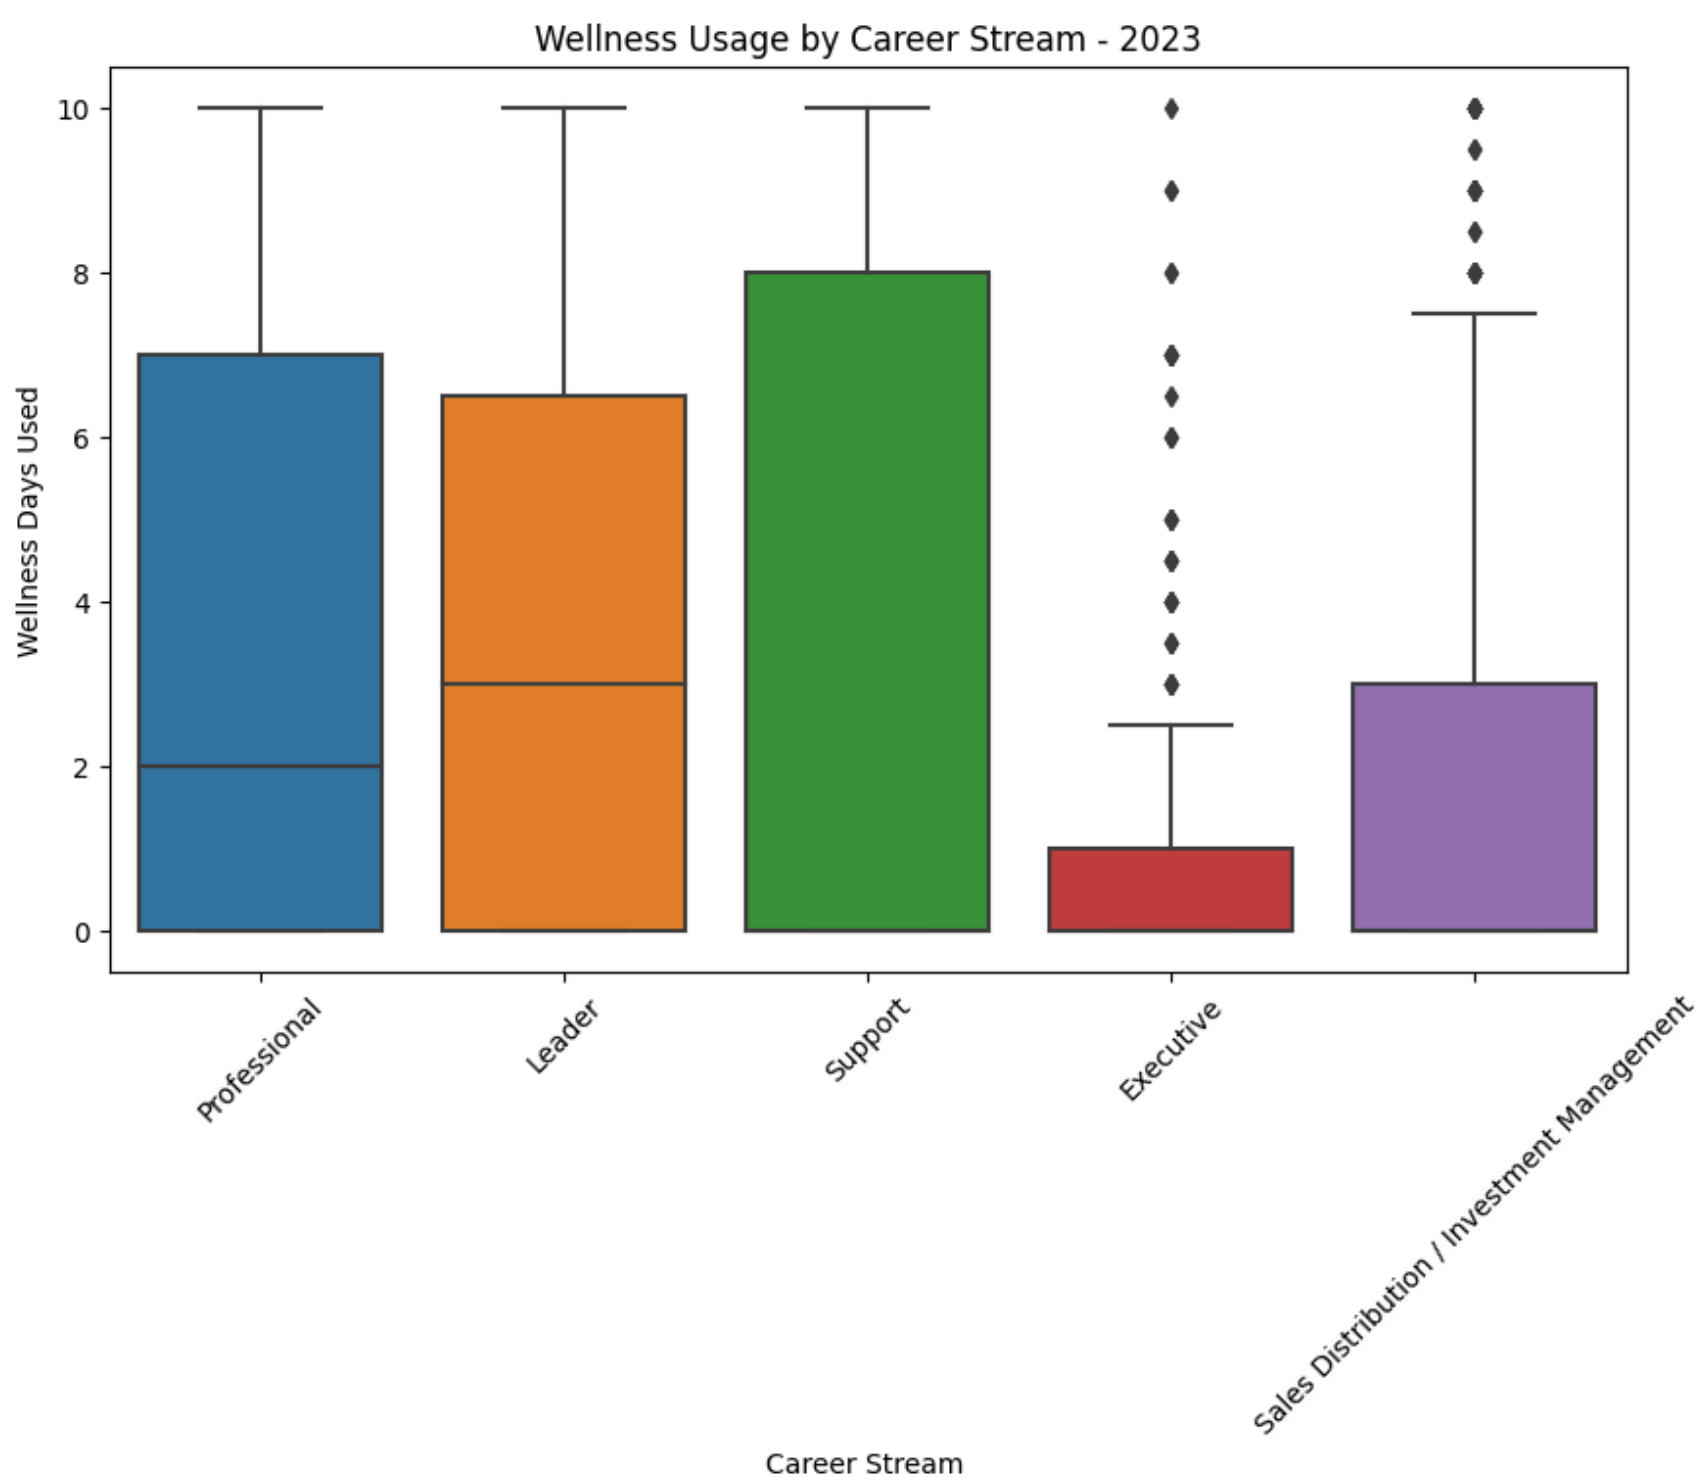
## WELLNESS USAGE BY CAREER STREAM

```
In [16]: # Plot box plot of wellness usage by Career stream for 2022
plt.figure(figsize=(10, 6))
sns.boxplot(x='Career Stream_x', y='Wellness - Used_x', data=merged_df)
plt.title('Wellness Usage by career Stream - 2022')
plt.xlabel('Career Stream')
plt.ylabel('Wellness Days Used')
plt.xticks(rotation=45)
plt.show()
```



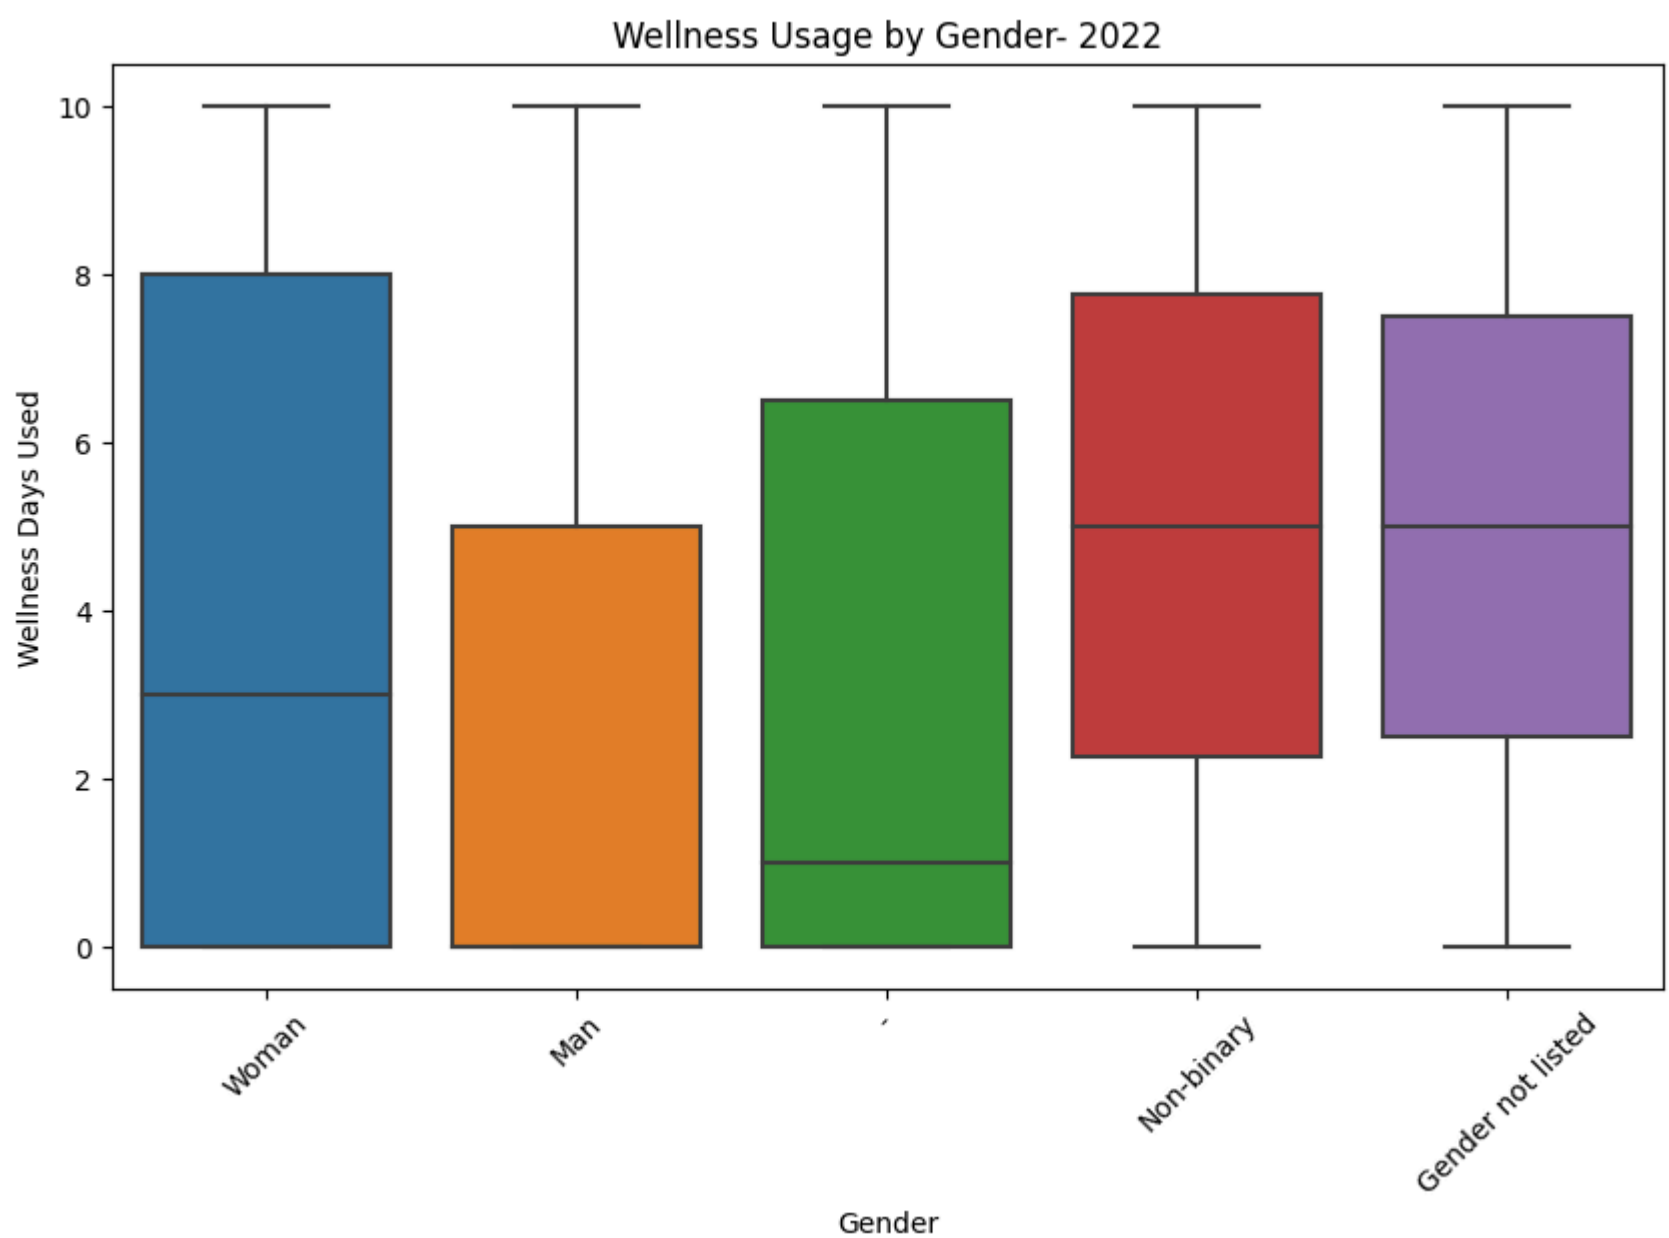


```
In [17]: # Plot box plot of wellness usage by Career stream for 2023
plt.figure(figsize=(10, 6))
sns.boxplot(x='Career Stream_y', y='Wellness - Used_y', data=merged_df)
plt.title('Wellness Usage by Career Stream - 2023')
plt.xlabel('Career Stream ')
plt.ylabel('Wellness Days Used')
plt.xticks(rotation=45)
plt.show()
```




## WELLNESS USAGE BY GENDER

```
In [18]: # Plot box plot of wellness usage by Gender for 2022
plt.figure(figsize=(10, 6))
sns.boxplot(x='Gender_x', y='Wellness - Used_x', data=merged_df)
plt.title('Wellness Usage by Gender- 2022')
plt.xlabel('Gender')
plt.ylabel('Wellness Days Used')
plt.xticks(rotation=45)
plt.show()
```




```
In [19]: # Plot box plot of wellness usage by Gender for 2023
plt.figure(figsize=(10, 6))
sns.boxplot(x='Gender_y', y='Wellness - Used_y', data=merged_df)
plt.title('Wellness Usage by Gender - 2023')
plt.xlabel('Gender')
plt.ylabel('Wellness Days Used')
plt.xticks(rotation=45)
plt.show()
```



**CORRELATION BETWEEN WELLNESS USED, AGE, AND VACATION USED**

```
In [20]: # Calculate correlations between wellness usage and other numeric variables
wellness_corr = merged_df[['Wellness - Used_x', 'Wellness - Used_y', 'Age_x','Age_y', 'Vacation - Used_x', 'Vacation - Used_y']]

# Plot correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(wellness_corr, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix: Wellness vs. Other Variables')
plt.show()
```



## Hypothesis Testing: Wellness Usage and Employee Retention

### Hypothesis

- **Null Hypothesis (H0):** There is no significant difference in wellness usage between employees who stay and employees who leave the company.
- **Alternative Hypothesis (H1):** Employees who stay use significantly more wellness benefits compared to employees who leave.

### Methodology

- I performed a two-sample t-test to compare the mean wellness usage between employees who are active (stay) and those who are terminated (leave).
- The significance level (alpha) chosen was 0.05.

```
In [21]: import scipy.stats as stats

# Separate wellness usage data for employees who stay and those who leave
wellness_stay = merged_df.loc[merged_df['Employee Status_x'] == 'Active', 'Wellness - Used_x']
wellness_leave = merged_df.loc[merged_df['Employee Status_x'] == 'Terminated', 'Wellness - Used_x']

In [22]: # Perform two-sample t-test
t_stat, p_value = stats.ttest_ind(wellness_stay, wellness_leave, equal_var=True)

In [23]: # Print t-statistic and p-value
print("T-statistic:", t_stat)
print("P-value:", p_value)

T-statistic: 59.54929608933553
P-value: 0.0
```

```
In [24]: # Interpret the results
alpha = 0.05 # significance level
if p_value < alpha:
    print("Reject null hypothesis: There is a significant difference in wellness usage between employees who stay and those who leave.")
else:
    print("Fail to reject null hypothesis: There is no significant difference in wellness usage between employees who stay and those who leave.")
```

Reject null hypothesis: There is a significant difference in wellness usage between employees who stay and those who leave.

## Results

- **T-Statistic:** 59.55
- **P-Value:** 0.0 (rounded)

## Conclusion

- Based on the results of the t-test:
  - **Reject Null Hypothesis:** There is a significant difference in wellness usage between employees who stay and those who leave the company.
  - Employees who stay tend to use significantly more wellness benefits compared to employees who leave.


## Correlation Analysis:

I'll start by calculating the correlation coefficients between wellness usage and other relevant variables such as age, vacation days used, etc. This will help us understand the strength and direction of these relationships.

```
In [25]: # Calculate correlation matrix
correlation_matrix = merged_df[['Wellness - Used_x', 'Age_x', 'Vacation - Used_x']].corr()

# Plot correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix: Wellness vs. Age vs. Vacation Days Used')
plt.show()

#Print correlation matrix
print("Correlation Matrix:")
print(corr_matrix)
```



Correlation Matrix:

	Employee ID	Vacation Annual Entitlement_x \
Employee ID	1.000000	-0.602802
Vacation Annual Entitlement_x	-0.602802	1.000000
Age_x	-0.350633	0.534915
Wellness - Used_x	-0.086591	0.082760
Wellness - Accrued_x	-0.099079	0.332575
Wellness - %	-0.079062	0.077978
Vacation - Used_x	-0.344016	0.502085
Vacation - Accrued_x	-0.336177	0.553775
Vacation - %	-0.180767	0.329327
Age_y	-0.378433	0.534970
Vacation Annual Entitlement_y	-0.617874	0.953540
Wellness - Used_y	-0.041682	0.035086
Wellness - Accrued_y	-0.049798	0.269746
Wellness Used- %	-0.033185	0.035113
Vacation - Used_y	-0.265518	0.387168
Vacation - Accrued_y	-0.245354	0.417957
Vacation Used - %	-0.105911	0.247412

	Age_x	Wellness - Used_x \
Employee ID	-0.350633	-0.086591
Vacation Annual Entitlement_x	0.534915	0.082760
Age_x	1.000000	0.062610
Wellness - Used_x	0.062610	1.000000
Wellness - Accrued_x	0.196006	0.631080
Wellness - %	0.061794	0.994040
Vacation - Used_x	0.294099	0.620251
Vacation - Accrued_x	0.316252	0.572308
Vacation - %	0.207012	0.685026
Age_y	0.999999	0.064155
Vacation Annual Entitlement_y	0.519121	0.132169
Wellness - Used_y	0.036885	0.697891
Wellness - Accrued_y	0.160822	0.511708
Wellness Used- %	0.036685	0.697990
Vacation - Used_y	0.221896	0.500125
Vacation - Accrued_y	0.232501	0.471604
Vacation Used - %	0.153658	0.532483

	Wellness - Accrued_x	Wellness - % \
Employee ID	-0.099079	-0.079062
Vacation Annual Entitlement_x	0.332575	0.077978
Age_x	0.196006	0.061794
Wellness - Used_x	0.631080	0.994040
Wellness - Accrued_x	1.000000	0.626796
Wellness - %	0.626796	1.000000
Vacation - Used_x	0.789010	0.612503
Vacation - Accrued_x	0.864607	0.565545
Vacation - %	0.841832	0.677569
Age_y	0.195404	0.063338
Vacation Annual Entitlement_y	0.381552	0.127193
Wellness - Used_y	0.572411	0.699353
Wellness - Accrued_y	0.786154	0.513710
Wellness Used- %	0.572475	0.699450
Vacation - Used_y	0.729010	0.499262
Vacation - Accrued_y	0.767673	0.471552
Vacation Used - %	0.737329	0.533616

	Vacation - Used_x	Vacation - Accrued_x \
Employee ID	-0.344016	-0.336177
Vacation Annual Entitlement_x	0.502085	0.553775
Age_x	0.294099	0.316252
Wellness - Used_x	0.620251	0.572308
Wellness - Accrued_x	0.789010	0.864607
Wellness - %	0.612503	0.565545
Vacation - Used_x	1.000000	0.921932
Vacation - Accrued_x	0.921932	1.000000
Vacation - %	0.930321	0.847012
Age_y	0.295451	0.316811
Vacation Annual Entitlement_y	0.530218	0.584768
Wellness - Used_y	0.485306	0.457552
Wellness - Accrued_y	0.655662	0.704196
Wellness Used- %	0.485380	0.457631
Vacation - Used_y	0.740302	0.737149
Vacation - Accrued_y	0.715766	0.773807
Vacation Used - %	0.675164	0.672772

	Vacation - %	Age_y \
Employee ID	-0.180767	-0.378433
Vacation Annual Entitlement_x	0.329327	0.534970
Age_x	0.207012	0.999999
Wellness - Used_x	0.685026	0.064155
Wellness - Accrued_x	0.841832	0.195404
Wellness - %	0.677569	0.063338
Vacation - Used_x	0.930321	0.295451
Vacation - Accrued_x	0.847012	0.316811
Vacation - %	1.000000	0.209061
Age_y	0.209061	1.000000
Vacation Annual Entitlement_y	0.376393	0.535061
Wellness - Used_y	0.552744	0.055416
Wellness - Accrued_y	0.688758	0.152710
Wellness Used- %	0.552821	0.054413
Vacation - Used_y	0.717371	0.249812
Vacation - Accrued_y	0.691663	0.260653
Vacation Used - %	0.706062	0.157626

	Vacation Annual Entitlement_y \	
Employee ID	-0.617874	
Vacation Annual Entitlement_x	0.953540	
Age_x	0.519121	
Wellness - Used_x	0.132169	
Wellness - Accrued_x	0.381552	
Wellness - %	0.127193	
Vacation - Used_x	0.530218	
Vacation - Accrued_x	0.584768	
Vacation - %	0.376393	
Age_y	0.535061	
Vacation Annual Entitlement_y	1.000000	
Wellness - Used_y	0.083247	
Wellness - Accrued_y	0.277843	
Wellness Used- %	0.077647	
Vacation - Used_y	0.446862	
Vacation - Accrued_y	0.476740	
Vacation Used - %	0.270857	

	Wellness - Used_y	Wellness - Accrued_y \
Employee ID	-0.041682	-0.049798
Vacation Annual Entitlement_x	0.035086	0.269746
Age_x	0.036885	0.160822
Wellness - Used_x	0.697891	0.511708
Wellness - Accrued_x	0.572411	0.786154
Wellness - %	0.699353	0.513710
Vacation - Used_x	0.485306	0.655662
Vacation - Accrued_x	0.457552	0.704196
Vacation - %	0.552744	0.688758
Age_y	0.055416	0.152710
Vacation Annual Entitlement_y	0.083247	0.277843
Wellness - Used_y	1.000000	0.681465
Wellness - Accrued_y	0.681465	1.000000
Wellness Used- %	0.994161	0.673699
Vacation - Used_y	0.656795	0.822174
Vacation - Accrued_y	0.631674	0.885636
Vacation Used - %	0.703030	0.859778

	Wellness Used- %	Vacation - Used_y \
Employee ID	-0.033185	-0.265518
Vacation Annual Entitlement_x	0.035113	0.387168
Age_x	0.036685	0.221896
Wellness - Used_x	0.697990	0.500125
Wellness - Accrued_x	0.572475	0.729010
Wellness - %	0.699450	0.499262
Vacation - Used_x	0.485380	0.740302
Vacation - Accrued_x	0.457631	0.737149
Vacation - %	0.552821	0.717371
Age_y	0.054413	0.249812
Vacation Annual Entitlement_y	0.077647	0.446862
Wellness - Used_y	0.994161	0.656795
Wellness - Accrued_y	0.673699	0.822174
Wellness Used- %	1.000000	0.648686
Vacation - Used_y	0.648686	1.000000
Vacation - Accrued_y	0.624710	0.941791
Vacation Used - %	0.699995	0.933705

	Vacation - Accrued_y	Vacation Used - %
Employee ID	-0.245354	-0.105911
Vacation Annual Entitlement_x	0.417957	0.247412
Age_x	0.232501	0.153658
Wellness - Used_x	0.471604	0.532483
Wellness - Accrued_x	0.767673	0.737329
Wellness - %	0.471552	0.533616
Vacation - Used_x	0.715766	0.675164
Vacation - Accrued_x	0.773807	0.672772
Vacation - %	0.691663	0.706062
Age_y	0.260653	0.157626
Vacation Annual Entitlement_y	0.476740	0.270857
Wellness - Used_y	0.631674	0.703030
Wellness - Accrued_y	0.885636	0.859778
Wellness Used- %	0.624710	0.699995
Vacation - Used_y	0.941791	0.933705
Vacation - Accrued_y	1.000000	0.870477
Vacation Used - %	0.870477	1.000000

## Regression Analysis:

I perform a regression analysis to model the relationship between wellness usage (independent variable) and employee retention (dependent variable). Since employee retention is a binary outcome (stay or leave), I use logistic regression.

```
In [26]: import pandas as pd
import statsmodels.api as sm

# Drop rows with missing values in relevant columns
merged_df_cleaned = merged_df.dropna(subset=['Wellness - Used_x', 'Employee Status_x'])

# Define predictors (X) and target (y) after handling missing values
X = merged_df_cleaned[['Wellness - Used_x']]
y = (merged_df_cleaned['Employee Status_x'] == 'Active').astype(int) # Convert 'Active' to 1, 'Terminated' to 0
```

```
# Add constant to X for intercept in regression
X = sm.add_constant(X)

# Fit Logistic regression model
logit_model = sm.Logit(y, X)
result = logit_model.fit()

# Print model summary
print(result.summary())
```

Optimization terminated successfully.  
Current function value: 0.419404  
Iterations 8

Logit Regression Results						
=====						
Dep. Variable:	Employee Status_x	No. Observations:	6044			
Model:	Logit	Df Residuals:	6042			
Method:	MLE	Df Model:	1			
Date:	Thu, 13 Jun 2024	Pseudo R-squ.:	0.3533			
Time:	11:27:33	Log-Likelihood:	-2534.9			
converged:	True	LL-Null:	-3919.7			
Covariance Type:	nonrobust	LLR p-value:	0.000			
=====						
	coef	std err	z	P> z	[0.025	0.975]
-----						
const	-0.6206	0.038	-16.327	0.000	-0.695	-0.546
Wellness - Used_x	0.6801	0.023	30.121	0.000	0.636	0.724
=====						

## SURVIVAL ANALYSIS \*\*

# setting Survival Analysis: Time to Employee Termination

### Approach

I conducted survival analysis to understand the duration until an employee terminates employment using the Kaplan-Meier estimator.

### Methodology

1. **Data Preparation:**

- I calculated the tenure (time from hire to termination) in days using the 'Hire Date' and 'Termination Date' columns.
- Rows with missing tenure values were dropped to ensure data integrity.

2. **Survival Analysis:**

- I used the KaplanMeierFitter from the lifelines library to estimate the survival function.
- The analysis was based on time-to-event data (tenure) and event status (termination).

3. **Visualization:**

- A Kaplan-Meier survival curve was plotted to visualize the probability of employees not leaving the company over time.

### Results

The survival curve provides insights into the duration employees stay with the company before termination. It helps in understanding the attrition pattern and can guide HR strategies for employee retention and management.

```
In [27]: !pip install lifelines
```



Collecting lifelines  
 Downloading lifelines-0.28.0-py3-none-any.whl.metadata (3.2 kB)  
Requirement already satisfied: numpy<2.0,>=1.14.0 in /opt/conda/lib/python3.10/site-packages (from lifelines) (1.26.4)  
Requirement already satisfied: scipy>=1.2.0 in /opt/conda/lib/python3.10/site-packages (from lifelines) (1.11.4)  
Requirement already satisfied: pandas>=1.2.0 in /opt/conda/lib/python3.10/site-packages (from lifelines) (2.2.2)  
Requirement already satisfied: matplotlib>=3.0 in /opt/conda/lib/python3.10/site-packages (from lifelines) (3.7.5)  
Collecting autograd>=1.5 (from lifelines)  
 Downloading autograd-1.6.2-py3-none-any.whl.metadata (706 bytes)  
Collecting autograd-gamma>=0.3 (from lifelines)  
 Downloading autograd-gamma-0.5.0.tar.gz (4.0 kB)  
 Preparing metadata (setup.py) ... done  
Collecting formulaic>=0.2.2 (from lifelines)  
 Downloading formulaic-1.0.1-py3-none-any.whl.metadata (6.1 kB)  
Requirement already satisfied: future>=0.15.2 in /opt/conda/lib/python3.10/site-packages (from autograd>=1.5->lifelines) (1.0.0)  
Collecting interface-meta>=1.2.0 (from formulaic>=0.2.2->lifelines)  
 Downloading interface\_meta-1.3.0-py3-none-any.whl.metadata (6.7 kB)  
Requirement already satisfied: typing-extensions>=4.2.0 in /opt/conda/lib/python3.10/site-packages (from formulaic>=0.2.2->lifelines) (4.9.0)  
Requirement already satisfied: wrapt>=1.0 in /opt/conda/lib/python3.10/site-packages (from formulaic>=0.2.2->lifelines) (1.14.1)  
Requirement already satisfied: contourpy>=1.0.1 in /opt/conda/lib/python3.10/site-packages (from matplotlib>=3.0->lifelines) (1.2.0)  
Requirement already satisfied: cycler>=0.10 in /opt/conda/lib/python3.10/site-packages (from matplotlib>=3.0->lifelines) (0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in /opt/conda/lib/python3.10/site-packages (from matplotlib>=3.0->lifelines) (4.47.0)  
Requirement already satisfied: kiwisolver>=1.0.1 in /opt/conda/lib/python3.10/site-packages (from matplotlib>=3.0->lifelines) (1.4.5)  
Requirement already satisfied: packaging>=20.0 in /opt/conda/lib/python3.10/site-packages (from matplotlib>=3.0->lifelines) (21.3)  
Requirement already satisfied: pillow>=6.2.0 in /opt/conda/lib/python3.10/site-packages (from matplotlib>=3.0->lifelines) (9.5.0)  
Requirement already satisfied: pyparsing>=2.3.1 in /opt/conda/lib/python3.10/site-packages (from matplotlib>=3.0->lifelines) (3.1.1)  
Requirement already satisfied: python-dateutil>=2.7 in /opt/conda/lib/python3.10/site-packages (from matplotlib>=3.0->lifelines) (2.9.0.post0)  
Requirement already satisfied: pytz>=2020.1 in /opt/conda/lib/python3.10/site-packages (from pandas>=1.2.0->lifelines) (2023.3.post1)  
Requirement already satisfied: tzdata>=2022.7 in /opt/conda/lib/python3.10/site-packages (from pandas>=1.2.0->lifelines) (2023.4)  
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.10/site-packages (from python-dateutil>=2.7->matplotlib>=3.0->lifelines) (1.16.0)  
Downloading lifelines-0.28.0-py3-none-any.whl (349 kB)  
 349.2/349.2 kB 10.3 MB/s eta 0:00:00  
Downloading autograd-1.6.2-py3-none-any.whl (49 kB)  
 49.3/49.3 kB 2.5 MB/s eta 0:00:00  
Downloading formulaic-1.0.1-py3-none-any.whl (94 kB)  
 94.2/94.2 kB 4.1 MB/s eta 0:00:00  
Downloading interface\_meta-1.3.0-py3-none-any.whl (14 kB)  
Building wheels for collected packages: autograd-gamma  
 Building wheel for autograd-gamma (setup.py) ... done  
 Created wheel for autograd-gamma: filename=autograd\_gamma-0.5.0-py3-none-any.whl size=4031 sha256=319178f2c2339f6a100cfd5d8fcb55c36bd961a84085de1dfc6b593157049e58  
 Stored in directory: /root/.cache/pip/wheels/25/cc/e0/ef2969164144c899fedb22b338f6703e2b9cf46eeebf254991  
Successfully built autograd-gamma  
Installing collected packages: interface-meta, autograd, autograd-gamma, formulaic, lifelines  
Successfully installed autograd-1.6.2 autograd-gamma-0.5.0 formulaic-1.0.1 interface-meta-1.3.0 lifelines-0.28.0

```
In [28]: from lifelines import KaplanMeierFitter
import matplotlib.pyplot as plt
import pandas as pd

# Calculate tenure (time until termination) in days
merged_df['Tenure'] = (merged_df['Termination Date_x'] - merged_df['Hire Date_x']).dt.days


# Drop rows with missing tenure values (if any)
merged_df_cleaned = merged_df.dropna(subset=['Tenure', 'Employee Status_x'])

# Create a Kaplan-Meier estimator object
kmf = KaplanMeierFitter()

# Define time-to-event (survival time) and event status (1 for termination, 0 for ongoing employment)
time = merged_df_cleaned['Tenure']
event = (merged_df_cleaned['Employee Status_x'] == 'Terminated').astype(int)

# Fit the Kaplan-Meier estimator
kmf.fit(time, event)

# Plot the survival curve
plt.figure(figsize=(10, 6))
kmf.plot()
plt.title('Survival Curve: Time to Termination (using Hire Date)2022')
plt.xlabel('Time (Days)')
plt.ylabel('Survival Probability')
plt.show()
```



```
In [29]: # Calculate tenure (time from hire to termination) in days
merged_df['Tenure'] = (merged_df['Termination Date_y'] - merged_df['Hire Date_y']).dt.days

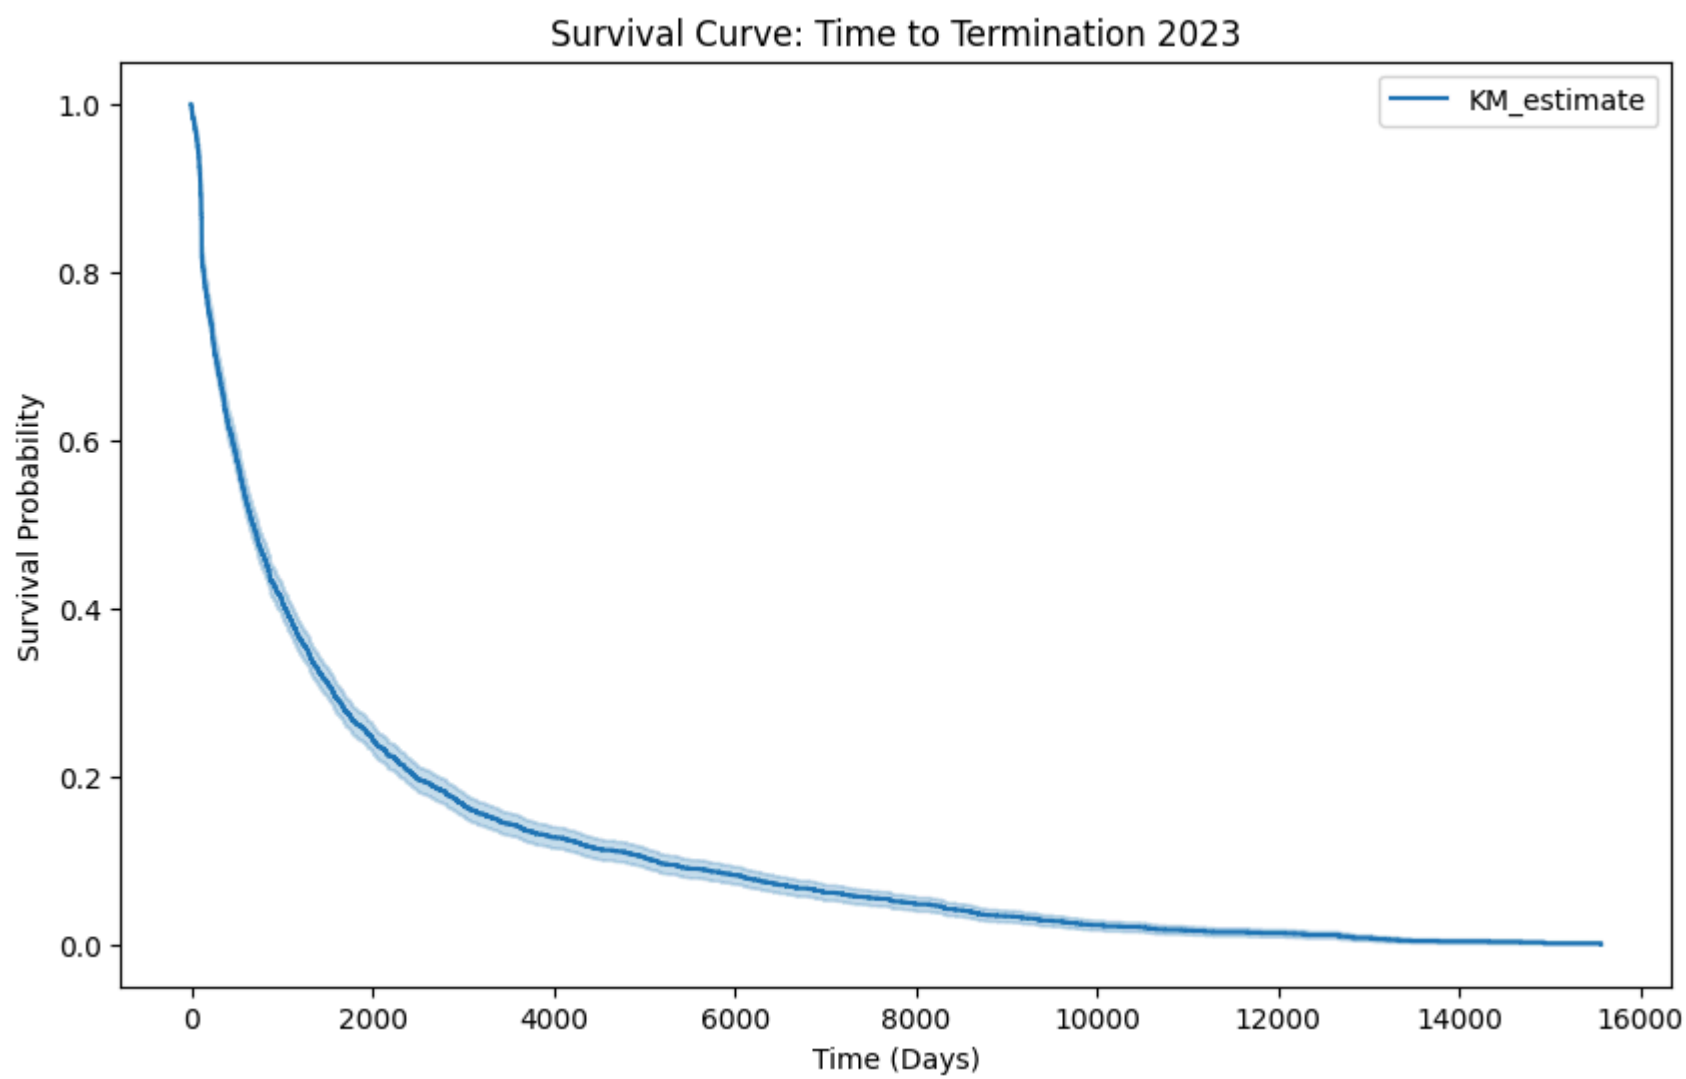
# Drop rows with missing tenure values (if any)
merged_df_cleaned = merged_df.dropna(subset=['Tenure', 'Employee Status_y'])

# Create a Kaplan-Meier estimator object
kmf = KaplanMeierFitter()

# Define time-to-event (survival time) and event status (1 for termination, 0 for ongoing employment)
time = merged_df_cleaned['Tenure']
event = (merged_df_cleaned['Employee Status_y'] == 'Terminated').astype(int)

# Fit the Kaplan-Meier estimator
kmf.fit(time, event)

# Plot the survival curve
plt.figure(figsize=(10, 6))
kmf.plot()
plt.title('Survival Curve: Time to Termination 2023')
plt.xlabel('Time (Days)')
plt.ylabel('Survival Probability')
plt.show()
```



# PREDICTIVE MODELING

## Predictive Modeling: Random Forest Classifier for Employee Retention

### Approach

I capabilities employed a Random Forest Classifier to predict the likelihood of an employee staying or leaving the company based on various features including wellness usage, age, and vacation days used.

### Methodology

1. **Data Preparation:**
- I selected 'Wellness - Used\_x', 'Age', and 'Vacation - Used\_x' as predictors (X) and 'Employee Status' as the target (y).
2. **Model Training and Testing:**
- The dataset was split into training and testing sets (80% training, 20% testing) using `train_test_split` from sklearn.
  - I trained the Random Forest Classifier with 100 estimators using `RandomForestClassifier` from sklearn.
3. **Model Evaluation:**
- I evaluated the model's performance on the test data using accuracy metrics.
  - The accuracy score, classification report, and confusion matrix were generated to assess the model's predictive capabilities.

```
In [30]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Drop rows with missing values in relevant columns
merged_df_cleaned = merged_df.dropna(subset=['Wellness - Used_x', 'Age_x', 'Vacation - Used_x', 'Employee Status_x'])

# Define predictors (X) and target (y) after handling missing values
X = merged_df_cleaned[['Wellness - Used_x', 'Age_x', 'Vacation - Used_x']]
y = (merged_df_cleaned['Employee Status_x'] == 'Active').astype(int) # Convert 'Active' to 1, 'Terminated' to 0

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Fit the model on training data
rf_classifier.fit(X_train, y_train)

# Make predictions on test data
y_pred = rf_classifier.predict(X_test)

# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate classification report and confusion matrix
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 0.9214226633581473

Classification Report:					
	precision	recall	f1-score	support	
0	0.88	0.90	0.89	426	
1	0.94	0.93	0.94	783	
accuracy			0.92	1209	
macro avg	0.91	0.92	0.91	1209	
weighted avg	0.92	0.92	0.92	1209	

Confusion Matrix:  
[[383 43]  
[ 52 731]]

## Results

- Accuracy:** 0.92 (rounded)
- Confusion Matrix:** ((383 43) (52 731))

### Conclusion

The Random Forest Classifier model provides insights into the factors influencing employee retention. Further analysis and feature engineering can enhance the model's predictive power and provide actionable insights for employee management strategies.

```
In [31]: import pandas as pd
df_1 = pd.read_excel('/kaggle/input/wellness2/2.xlsx', sheet_name = 'Diversity Data')

df_2 = pd.read_excel('/kaggle/input/wellness2/2.xlsx', sheet_name = '2023 Overtime Hours')

df_3 = pd.read_excel('/kaggle/input/wellness2/2.xlsx', sheet_name = '2022 Overtime Hours')

df_4 = pd.read_excel('/kaggle/input/wellness2/2.xlsx', sheet_name = '2022 Attrition Data')

df_5 = pd.read_excel('/kaggle/input/wellness2/2.xlsx', sheet_name = '2023 Attrition Data')

df_6 = pd.read_excel('/kaggle/input/wellness2/2.xlsx', sheet_name = '2023 Talent Review Data')

df_7 = pd.read_excel('/kaggle/input/wellness2/2.xlsx', sheet_name = '2023 Performance Data')


df_8 = pd.read_excel('/kaggle/input/wellness2/2.xlsx', sheet_name = '2022 Performance Data')
```

CORRELATION METRIX

```
In [32]: # Convert categorical data to numerical using one-hot encoding
df_encoded = pd.get_dummies(merged_df_cleaned[['Wellness - Used_x', 'Gender_x']])

corr_matrix_encoded = df_encoded.corr()

# Plot the correlation matrix as a heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix_encoded, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix of Wellness used and Gender')
plt.xlabel('Features')
plt.ylabel('Features')
plt.show()
```



```
In [33]: import seaborn as sns
import matplotlib.pyplot as plt

df_combined = pd.DataFrame()
df_combined['Wellness - Used_x'] = merged_df['Wellness - Used_x']
df_combined['Overall Performance Rating '] = df_8['Overall Performance Rating']

# Replace 'unrate' with NaN values in the DataFrame
df_combined.replace('Unrated', pd.NA, inplace=True)

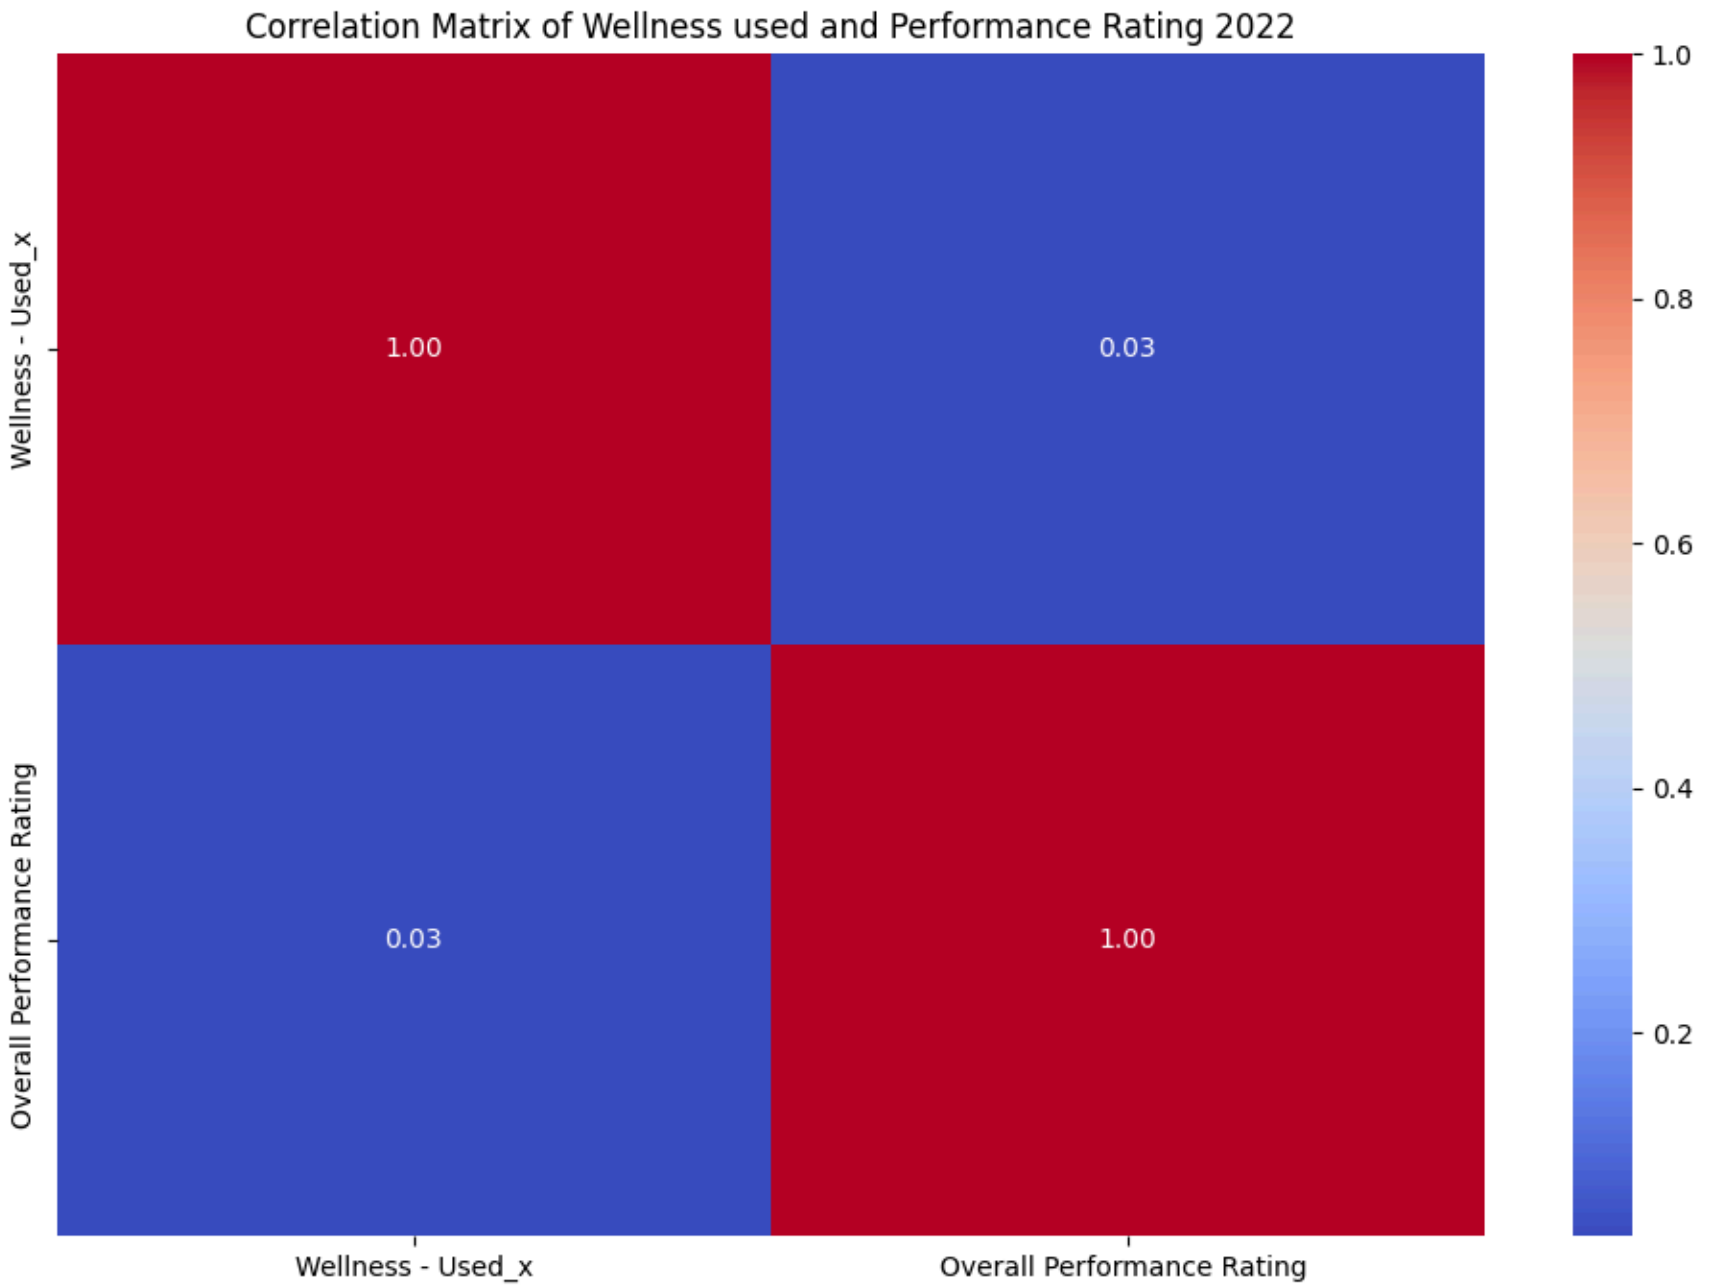
# Convert columns to numeric (if needed) to handle NaN values
df_combined = df_combined.apply(pd.to_numeric, errors='ignore')
```

```
# Drop rows with NaN values in any column
df_combined.dropna(inplace=True)

# Compute the correlation matrix for the combined DataFrame
corr_matrix_combined = df_combined.corr()

# Plot the correlation matrix as a heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix_combined, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix of Wellness used and Performance Rating 2022')
plt.show()

/tmp/ipykernel_18/467953103.py:15: FutureWarning: errors='ignore' is deprecated and will raise in a future version. Use
to_numeric without passing `errors` and catch exceptions explicitly instead
df_combined = df_combined.apply(pd.to_numeric, errors='ignore')
```



```
In [34]: import seaborn as sns
import matplotlib.pyplot as plt

df_combined = pd.DataFrame()
df_combined['Wellness - Used_y'] = merged_df['Wellness - Used_y']
df_combined['Overall Performance Rating'] = df_7['Overall Performance Rating']

# Replace 'unrate' with NaN values in the DataFrame
df_combined.replace('Unrated', pd.NA, inplace=True)

# Convert columns to numeric (if needed) to handle NaN values
df_combined = df_combined.apply(pd.to_numeric, errors='ignore')

# Drop rows with NaN values in any column
df_combined.dropna(inplace=True)

# Compute the correlation matrix for the combined DataFrame
corr_matrix_combined = df_combined.corr()

# Plot the correlation matrix as a heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix_combined, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix of Wellness used and Performance Rating 2023')
plt.show()

/tmp/ipykernel_18/1396843458.py:15: FutureWarning: errors='ignore' is deprecated and will raise in a future version. Use
e to_numeric without passing `errors` and catch exceptions explicitly instead
df_combined = df_combined.apply(pd.to_numeric, errors='ignore')
```

